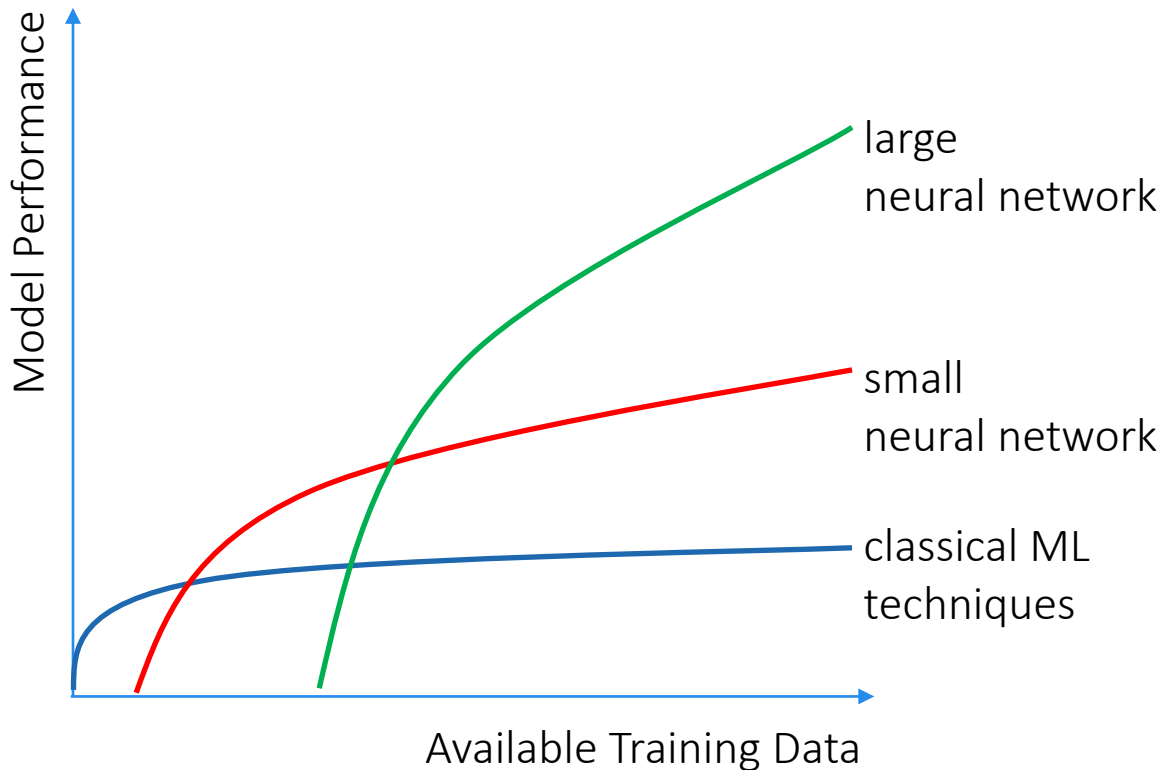


# Deep Learning Details

# Deep Learning Details

## Deep Learning Performance


- Classical ML techniques work best for small datasets
- With increasing size of available data → neural networks outperform classical techniques




# Deep Learning Details

Why did Deep Learning improve so much?


- Mainly four reasons why Deep Learning took off.



More  
Data



Moore's Law  
More computing  
Power  
GPU's



Better  
Algorithms





Open  
Source

# Deep Learning Details

Frameworks

S




 Microsoft  
CNTK theano

 mxnet

Caffe


# Deep Learning Details

## Perceptron




# Deep Learning Details

## Perceptron: Example






# Deep Learning Details

Simple and Deep Neural Network




Simple Neural Network

-  Input Values
-  Hidden Nodes
-  Output Nodes

# Deep Learning Details

Simple and Deep Neural Network



Deep Neural Network  
Multi-Layer Perceptron

- Input Values
- Hidden Nodes
- Output Nodes




# Deep Learning Details

## Layer Types

### Dense Layer

- Each input layer is connected to each output layer
- Also called fully connected layer
- Usually non-linear activation function applied




# Deep Learning Details

## Layer Types

### 1D convolutional layer

- Layer consists of filters
- Sequentially a subset of input layer is processed
- All nodes of input layer used




# Deep Learning Details

## Activation Functions

There are different activation functions.

### Rectified Linear Unit (ReLU)

- $\Phi = \max(0, x)$
- Most common
- Non-linear




# Deep Learning Details

## Activation Functions

### Leaky Rectified Linear Unit (Leaky ReLU)

- $\Phi(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha * x & \text{otherwise} \end{cases}$

- $\alpha$  typically 0.01
- Instead zero for negative inputs, small gradient
- Gradient never zero




# Deep Learning Details

## Activation Functions

### Hyperbolic Tangent (tanh)

- $\Phi(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
- Non-linear
- Relatively flat, except for small range
- Derivative small except for small range
- Might suffer vanishing gradient problem




# Deep Learning Details

## Activation Functions

### Sigmoid

- $\text{Phi}(x) = 1 / (1 + e^{-x})$
- Non-linear
- Relatively flat, except for small range
- Derivative small except for small range
- Might suffer vanishing gradient problem
- Result range 0 to 1

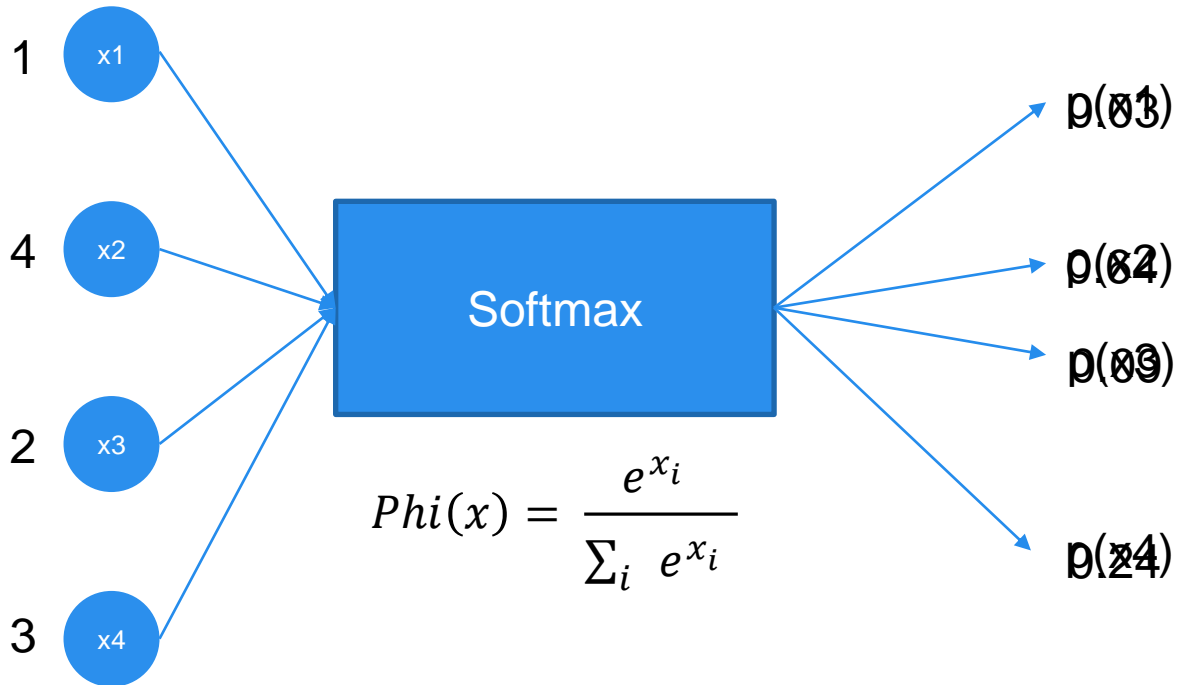


# Deep Learning Details

## Activation Functions

### Softmax

- Used for multi-class prediction



# Deep Learning Details

## Loss Functions

- Evaluates model performance during training
- Gradual improvement due to optimizer
- Is minimized during training
- Multiple loss functions for one model possible (one for each output variable)

### Regression

- Mean Squared Error  $MSE = \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{n}$
- Mean Absolute Error  $MAE = \frac{\sum_{i=1}^n |y_i - \hat{y}_i|}{n}$

### Classification

- Cross Entropy  
$$CE = -(y_i \log \hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)$$



# Deep Learning Details

## Optimizer

- Calculates updates of weights based on Loss Function
- implements gradient descent

# Deep Learning Details

## Wrapup

What you should know after this lecture:

- Simple Perceptron
- Multi-Layer Perceptron
- Most common layer types
- Most common activation functions
- Most common loss functions


# Deep Learning Details

## Tensor


- General representation of numeric objects




0-dimensional Tensor



1-dimensional Tensor



2-dimensional Tensor



3-dimensional Tensor