

Vector Databases

Vector Databases

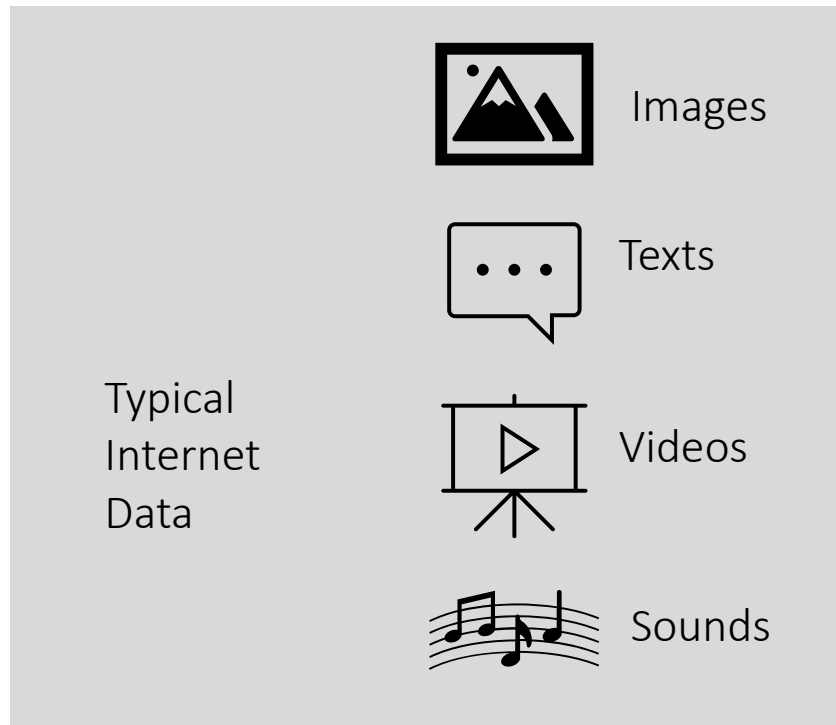
Introduction

- Special type of database
- Store, manage, and query data which is represented in geometric formats
- Very famous in Deep Learning (NLP)
- Allow
 - similarity search
 - Clustering
 - Real-time analytics


A vector database stores [high-dimensional data](#) (embeddings) for fast querying and similarity analysis.

Vector Databases


Why are Vector Databases required?



Unstructured Data



Typical
Databases





SQL

Structured Data

Vector Databases

Example: Image Similarity



Candidate Images




Query Image


Task: Which image is most similar?

Vector Databases

Example: Text Similarity



Candidate Words/Texts

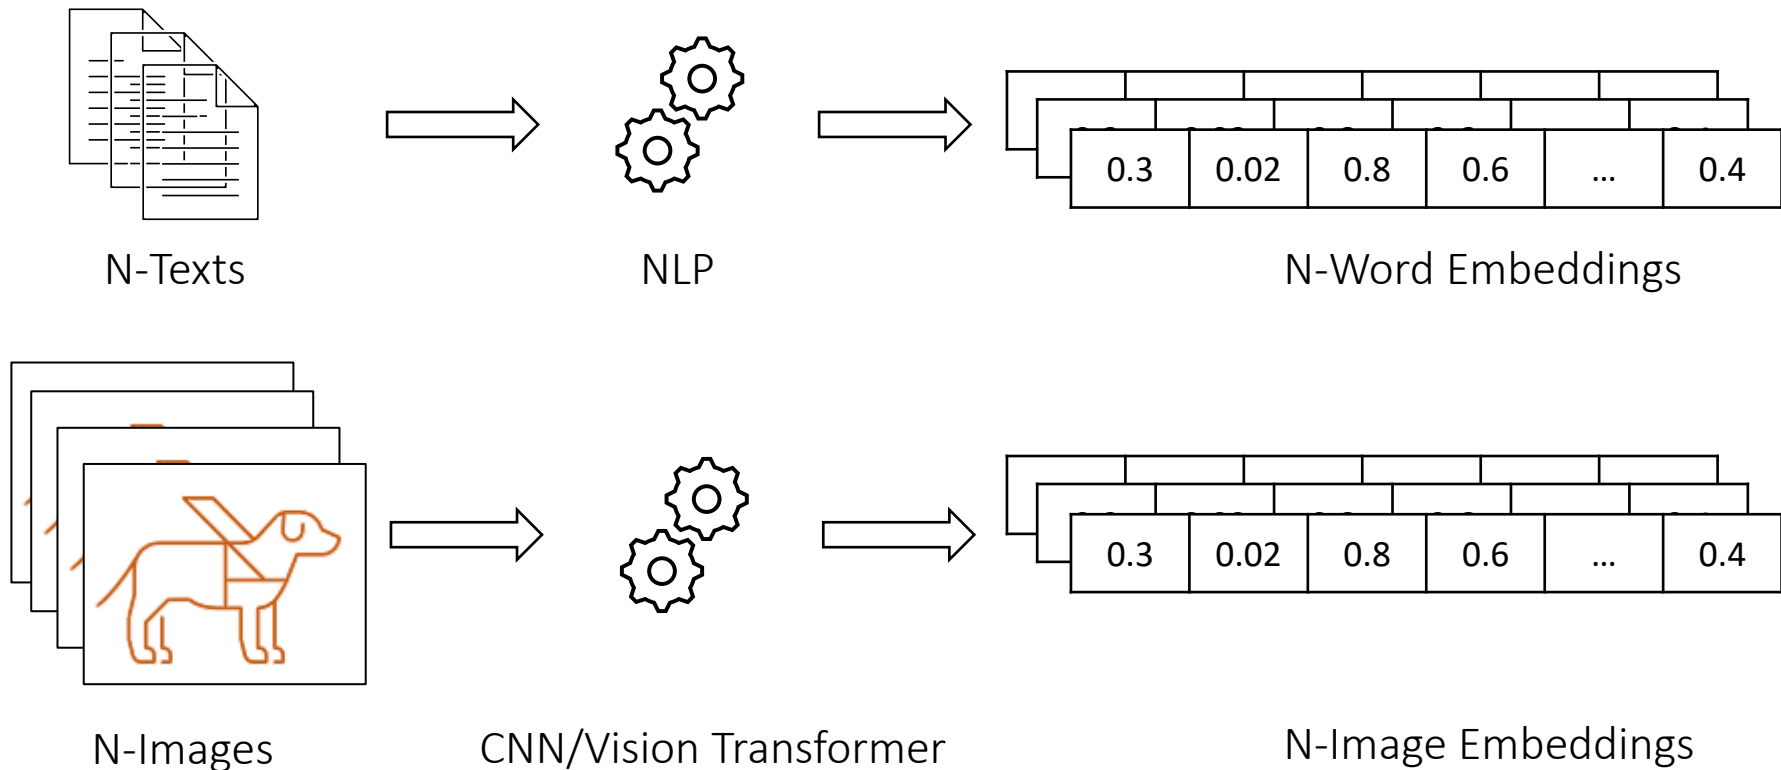


Query Words/Texts

Task: Which text is most similar?

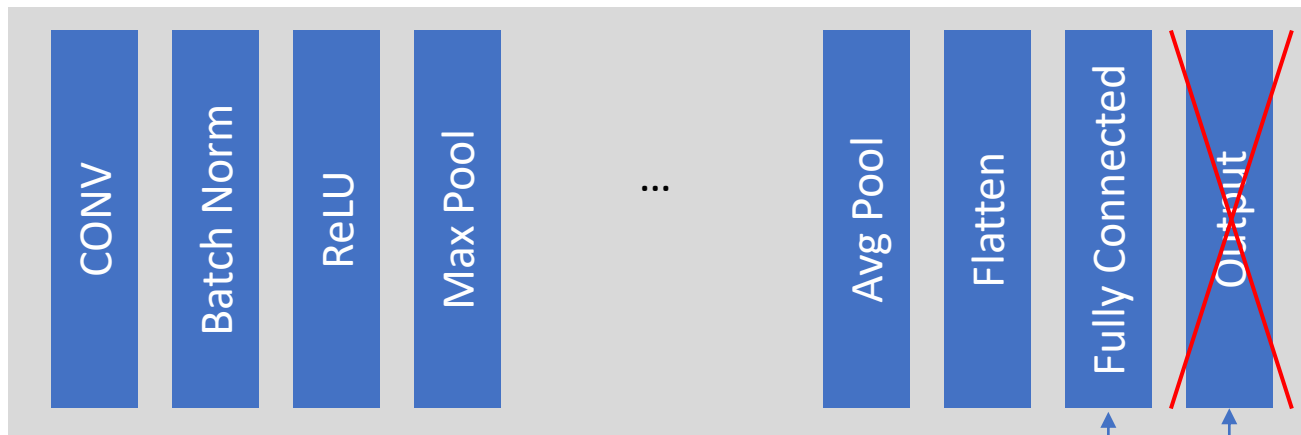

Vector Databases

Embeddings



Vector Databases

How to create image embeddings?




Last hidden state

Classes 0-999

Vector Databases

Find most similar data



$$dist = \sqrt{\underbrace{(x_1 - y_1)^2 + (x_n - y_n)^2}_{\text{distance terms}}}$$

For an embedding vector of 768 embeddings, there are 768 distance terms

Example: word embeddings reduced to 2 dimensions

Vector Databases

Storing embeddings / Calculate Distances

Imagename	Embedding					
dog1	0.3	0.02	0.8	0.6	...	0.4
dog2	0.1	0.52	0.7	0.6	...	0.4
...						
dogN	0.3	0.62	0.9	0.2	...	0.3

Vector Database

dogTest					
0.3	0.02	0.8	0.6	...	0.4






$$dist = \sqrt{\sum (x_i - y_i)^2}$$

For small data, go with np.array().

For Big Data this gets slow.

Vector Databases

Indexing

Imagename	Embedding						Index
dog1	0.3	0.02	0.8	0.6	...	0.4	
dog2	0.1	0.52	0.7	0.6	...	0.4	
...							
dogN	0.3	0.62	0.9	0.2	...	0.3	

Index speeds up search.

Vector Database

Vector Databases

Available Products



pinecone.io



trychroma.com



redis.com